

Predicting Twitch Channel Growth

A Regression Analysis of Streaming Activity, Audience Engagement, and Follower Growth

This report summarizes an end-to-end analysis of 1,000 top Twitch channels, aimed at understanding what drives channel growth and how well it can be predicted from streaming activity and audience metrics. The full workflow - data cleaning, exploratory analysis, feature engineering, model comparison, and overfitting diagnostics - is documented in two companion Jupyter notebooks; this report focuses on the key findings.

1,000

Channels Analyzed

11

Raw Features

21

Languages Represented

5

Regression Models
Compared

Growth metric used: We define channel growth as **Followers gained** - the number of new followers a channel gained over the observation period. This is the most direct, unambiguous growth signal available in the dataset (as opposed to point-in-time totals like current follower count).

1. Data Overview & Cleaning

The dataset covers 1,000 Twitch channels with watch time, stream time, peak/average concurrent viewers, follower counts and follower gains, views gained, and channel attributes (Partnered, Mature, Language). The data was found to be clean: no missing values and no duplicate rows.

One nuance surfaced during cleaning: three channels show a **negative** value for Followers gained - a net loss of followers over the period. Rather than treating this as an error, we treated it as legitimate churn signal and kept these rows, using a signed-log transform later so the model can still learn from them.

Almost every numeric column in the raw data is strongly right-skewed - a small number of mega-channels (e.g. the largest streamers) dominate the raw scale, while most channels cluster near the bottom. This is typical “power-law” platform data and directly motivated the log-transform approach used throughout feature engineering.

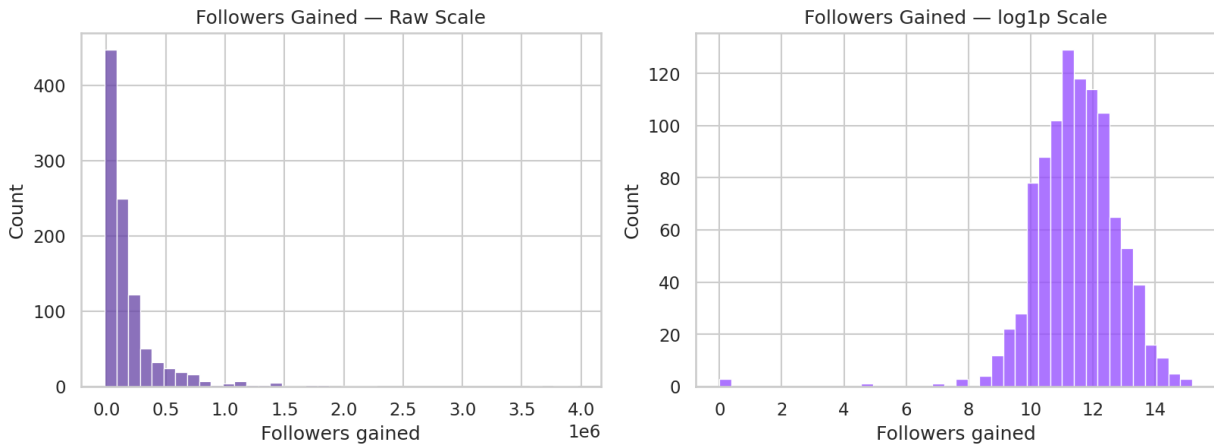


Figure 1. Followers Gained on raw scale (left) vs. $\log_{1p}$ scale (right). The log transform converts the heavy right skew into a much more symmetric, modelable distribution.

2. Exploratory Data Analysis - Key Relationships

On the log scale, several features correlate meaningfully with follower growth. The strongest relationships, in order, are: current Followers ($r \approx 0.72$), Watch time ($r \approx 0.51$), Peak viewers ($r \approx 0.47$), Average viewers ($r \approx 0.42$), and Views gained ($r \approx 0.24$).

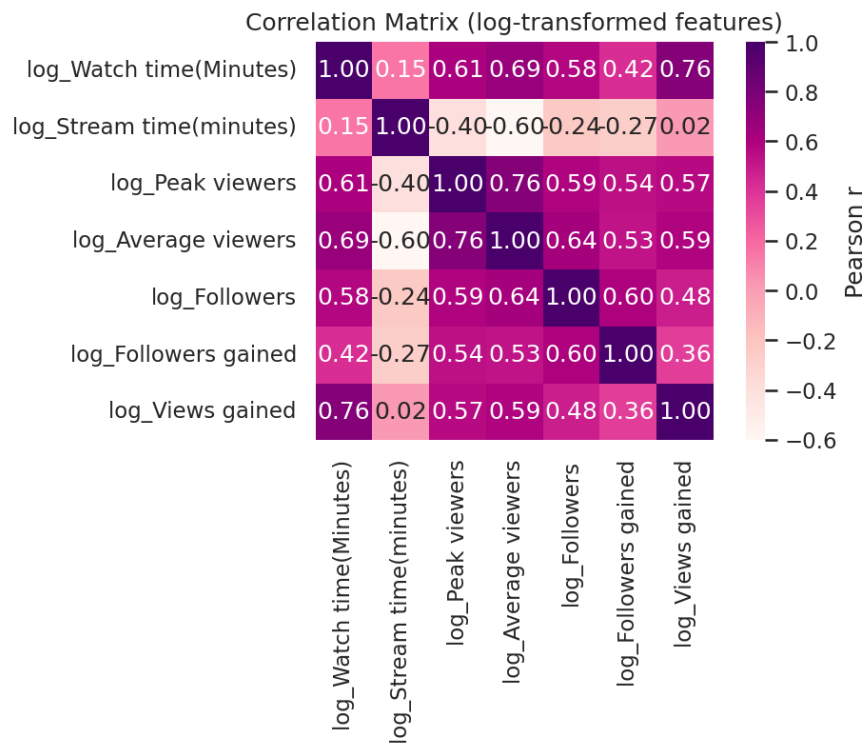


Figure 2. Correlation matrix of log-transformed numeric features.

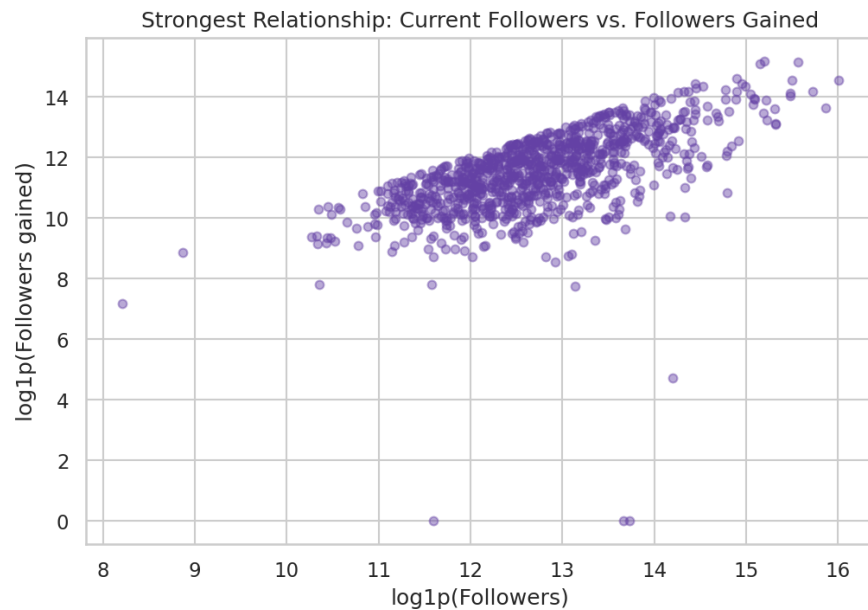


Figure 3. The single strongest bivariate relationship: a channel's current follower base vs. its follower gain (log-log scale).

Important caveat: current Followers is measured at the same snapshot as the growth target, so it partially overlaps with what it's predicting (a channel's current size already reflects the followers it just gained). It's an excellent predictor but a poor *actionable lever* - you can't grow your channel by "having more followers." We account for this directly in the modeling section.

Categorical attributes also show meaningful patterns. English dominates the channel count (about half of all channels), but growth varies noticeably by language. Only 2% of channels are non-Partnered, so that flag carries little statistical signal. Mature-flagged channels show a modestly different growth distribution than non-Mature channels.

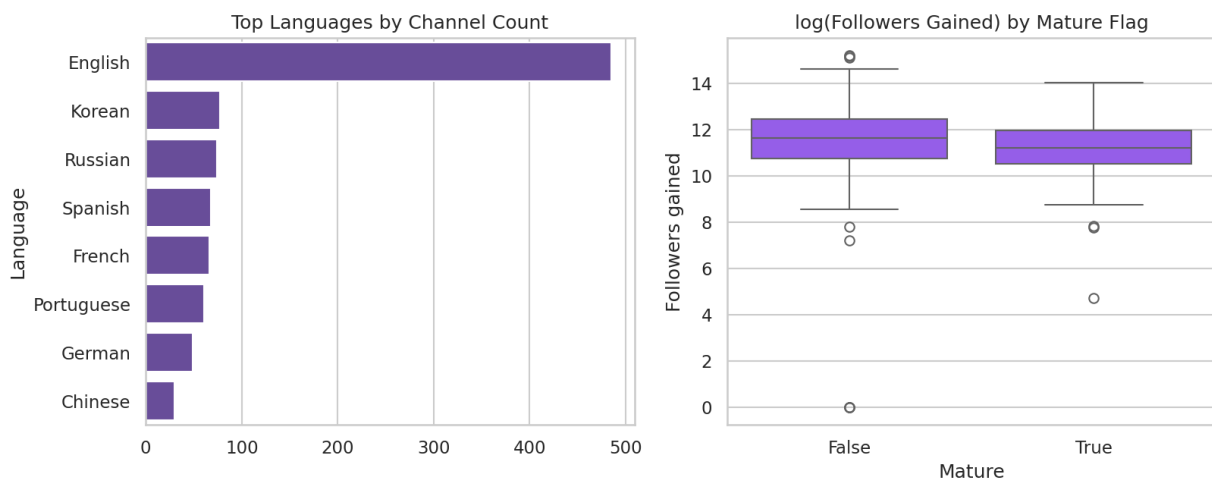


Figure 4. Left: channel counts by top languages. Right: $\log(\text{Followers gained})$ distribution split by Mature flag.

3. Feature Engineering

Building on the EDA findings, the following features were engineered for modeling:

- **Log-transformed predictors** - Watch time, Stream time, Peak viewers, Average viewers, Views gained, and Followers, to counter the heavy skew identified in EDA.
- **Engagement ratios** that normalize for channel size/activity: average-to-peak viewer ratio (audience consistency), watch-time-per-stream-minute (stickiness), and views-per-follower (discovery efficiency).
- **Categorical encoding** - Partnered and Mature as binary flags; Language grouped into the 6 most common languages plus an “Other” bucket, one-hot encoded (avoids an unwieldy number of sparse dummy columns for rare languages).
- **Signed-log growth target** - $\text{sign}(y) \cdot \log_{1p}(|y|)$ applied to Followers gained, which compresses the extreme right skew while still correctly handling the 3 channels with negative growth.

As noted above, we built two parallel feature sets to isolate the effect of the current-Followers leakage risk:

- **Set A** - includes current Followers (log-transformed). Best raw predictive power; serves as an upper-bound benchmark.
- **Set B** - activity-only (excludes current Followers). More honest measure of what streaming behavior actually drives growth, and the feature set we use for our primary, actionable model.

4. Model Comparison

Five regression approaches were trained and compared on both feature sets: Linear Regression, Ridge, Lasso, Random Forest, and Gradient Boosting. Each was evaluated with an 80/20 train-test split plus 5-fold cross-validation on the training data.

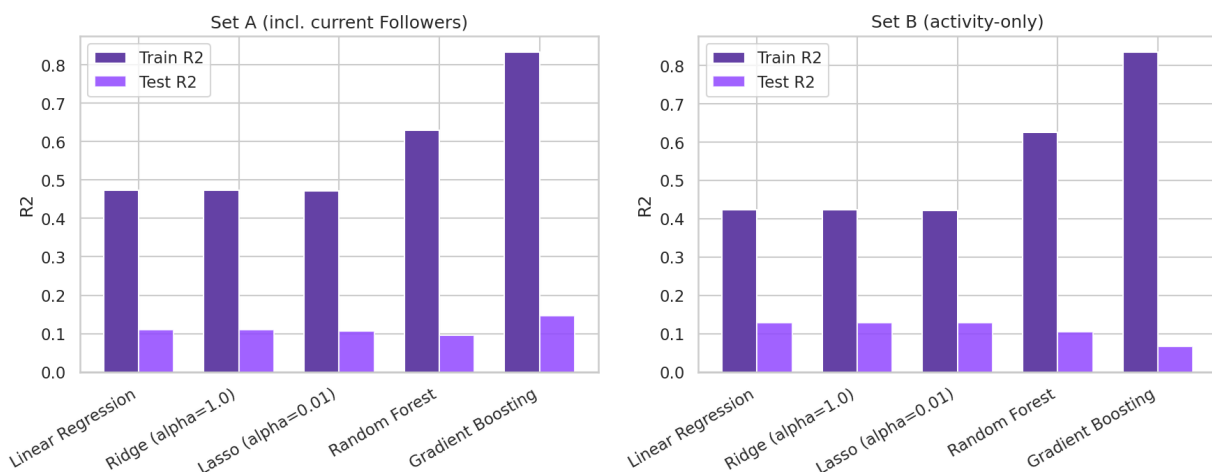


Figure 5. Train vs. Test R^2 for all five models, on both feature sets. Wide gaps between the two bars indicate overfitting.

Table 1. Set A results (includes current Followers)

Model	Train R^2	Test R^2	CV R^2	Overfit Gap
-------	-------------	------------	----------	-------------

Linear Regression	0.474	0.110	0.492	0.364
Ridge ($\alpha=1.0$)	0.474	0.110	0.492	0.364
Lasso ($\alpha=0.01$)	0.471	0.107	0.491	0.364
Random Forest	0.629	0.096	0.463	0.534
Gradient Boosting	0.833	0.146	0.350	0.687

Table 2. Set B results (activity-only, no leakage) - primary model set

Model	Train R ²	Test R ²	CV R ²	Overfit Gap
Linear Regression	0.424	0.129	0.438	0.295
Ridge ($\alpha=1.0$)	0.424	0.129	0.438	0.295
Lasso ($\alpha=0.01$)	0.421	0.129	0.437	0.293
Random Forest	0.625	0.105	0.457	0.520
Gradient Boosting	0.835	0.066	0.259	0.769

Selected model: Linear Regression on Set B. It achieves the best (tied) test R² (0.129) among activity-only models, with a smaller overfit gap than either ensemble method. On the raw (back-transformed) scale, this model's typical prediction error is about $\pm 107,000$ followers gained, against a dataset median of about 98,000 - useful for directional insight, but with meaningful uncertainty on individual channels.

5. Overfitting Diagnostics

Checking for overfitting was a core requirement of this analysis. Three checks were used:

- **Train vs. test R^2 gap** - the tree-based ensembles show by far the largest gaps (Random Forest: 0.52, Gradient Boosting: 0.77 on Set B), meaning they fit training data far better than they generalize - a textbook overfitting signature, even after applying depth/leaf-size regularization.
- **5-fold cross-validation** on the training data corroborates this: Gradient Boosting's cross-validated R^2 (0.26 on Set B) is far below its training R^2 (0.84), while linear models are far more consistent across folds (~ 0.44 CV R^2 vs. 0.42 train R^2).
- **Learning curves** for the selected model show training and cross-validation scores converging as training data grows, rather than diverging - the pattern expected from a well-generalizing model rather than one memorizing the training set.

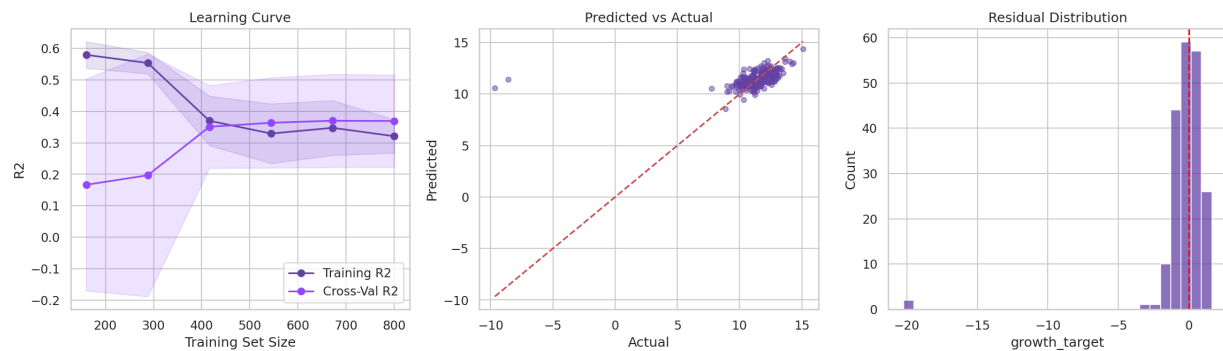


Figure 6. Left: learning curve for the final model (training and cross-validation scores converge with more data). Center: predicted vs. actual growth on the held-out test set. Right: residual distribution, centered near zero with no strong bias.

Honest read on absolute performance: even the best-generalizing model explains only a modest share of variance in follower growth (test $R^2 \approx 0.13$). This is a genuine finding, not a modeling failure to hide: channel growth on Twitch is driven by many factors this dataset doesn't capture (content quality, virality events, collaborations, platform algorithm changes, external promotion), so a snapshot of activity metrics alone should be expected to have real, meaningful limits in explaining it. The value of the model is in the *relative* signal it surfaces (Section 6), not in precise growth forecasting for any single channel.

6. What Actually Drives Growth?

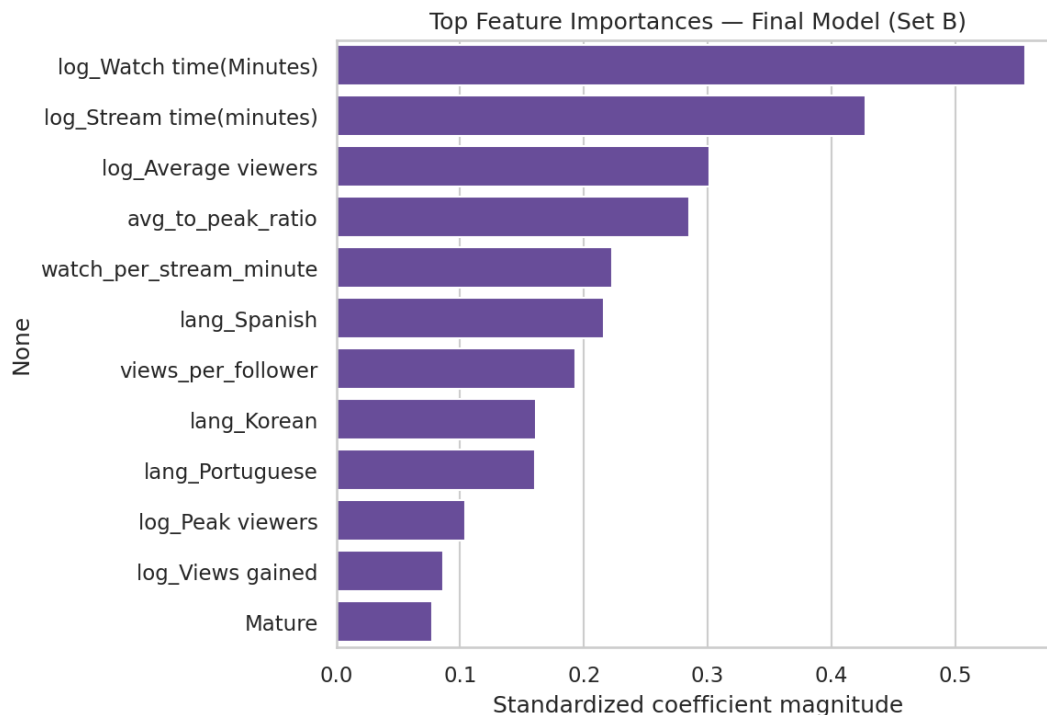


Figure 7. Standardized coefficient magnitudes for the final Linear Regression model (Set B) - larger bars indicate a stronger association with follower growth, holding other features constant.

The strongest activity-based drivers of follower growth, in order, are:

- **Watch time** - by far the top driver; channels that accumulate more total viewer-minutes gain more followers.
- **Stream time** - how much a channel streams overall is strongly associated with growth, independent of watch time.
- **Average viewers** and **average-to-peak viewer ratio** - channels with a consistently engaged audience (not just occasional viewership spikes) grow faster.
- **Watch-per-stream-minute** and **views-per-follower** - efficiency metrics that reward “sticky,” discoverable content over sheer volume.
- **Language** - Spanish, Korean, and Portuguese-language channels show a distinct growth association relative to the baseline group, suggesting regional/community dynamics matter.

By contrast, **Partnered status** contributes little (98% of channels are already Partnered, leaving little variation to learn from), and raw **Peak viewers** and **Views gained** matter less than the consistency- and efficiency-based ratios - reinforcing that steady engagement is a better growth signal than one-off spikes.

7. Conclusions & Recommendations

For streamers/creators: Consistency (watch time, stream time, and a high average-to-peak viewer ratio) is a more reliable growth lever than chasing viral peak-viewer moments alone. Building a steady, engaged base - reflected in average viewers and watch-per-minute stickiness - is associated with stronger follower growth than volume metrics like total views.

For platform/analytics teams: Predicting exact follower growth from activity metrics alone has real limits (test $R^2 \approx 0.13$) - growth is multi-causal, and a model like this is best used for directional prioritization (e.g., flagging channels with above-average watch-time/consistency for support or promotion) rather than precise forecasting.

Modeling takeaway: Simpler, regularized linear models generalized as well as or better than more complex tree-based ensembles on this dataset size (~1,000 rows), while being far less prone to overfitting. This is a useful reminder that model complexity should be matched to data volume - the more flexible Gradient Boosting model achieved the highest training fit but the worst test performance of all five models tested.

Limitations

- Single snapshot in time - no information on trend/trajectory prior to the measurement window.
- No content-level data (game/category, schedule regularity, collaborations, social media presence) that likely explain a meaningful share of the unexplained variance.
- Sample limited to the top 1,000 channels, so findings may not generalize to smaller or emerging channels with very different growth dynamics.

Full methodology, code, and additional diagnostics are available in the companion notebooks:

1_Data_Cleaning_EDA.ipynb and ***2_Feature_Engineering_Regression.ipynb***.